

# M5-Visual Recognition

## Week 1: PyTorch 101

### Group 07

Lali Bibilashvili  
Juan Chaves  
Carmen García  
Mohamed El Atiki



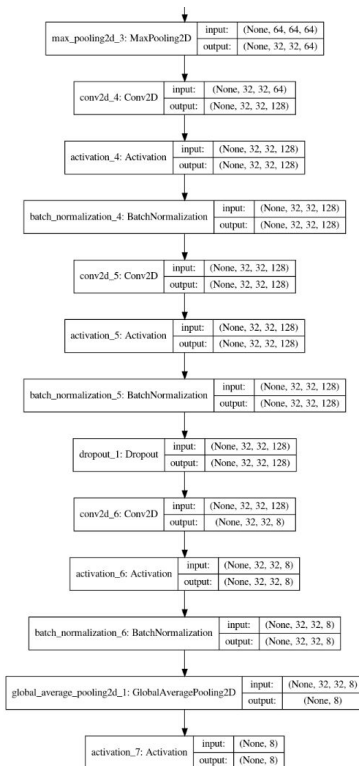
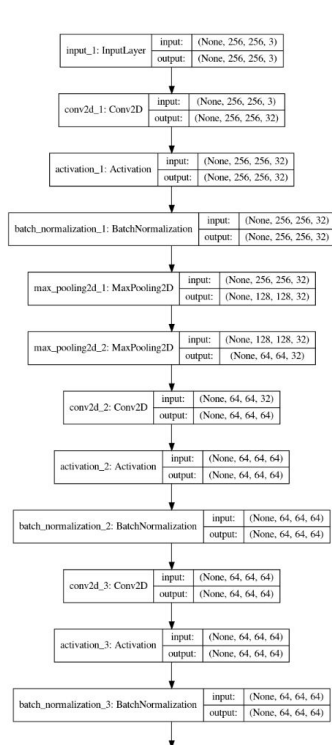
# Week 1

- This week focused on getting to know PyTorch by *translating* our last M3 Keras model
- As our group is formed by different M3 groups, we have worked on two models
- Apart from implementing the model itself, we also had to learn how to properly keep track of our performance and log our epochs statistics on Tensorboard
- We have also researched on some useful techniques, like Data Augmentation and Early Stopping



# First Model

# M3 - Model II



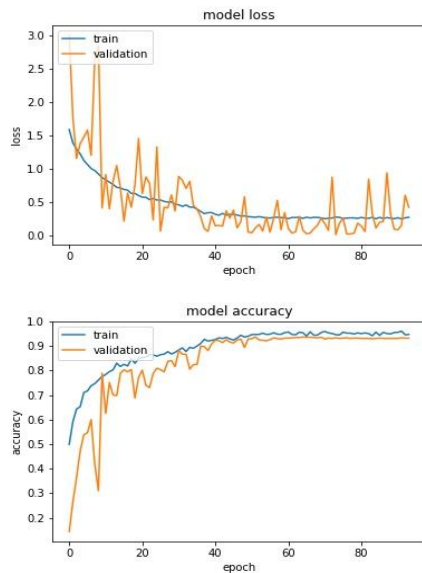
## Hyperparameters

Img. Size	256x256
Epochs	100
Batch Size	16
Learning Rate	1e-3


Activation	ReLu
Weight Init	Glorot uniform
Dropout rate	0.2

# M3 - Model II

## Keras



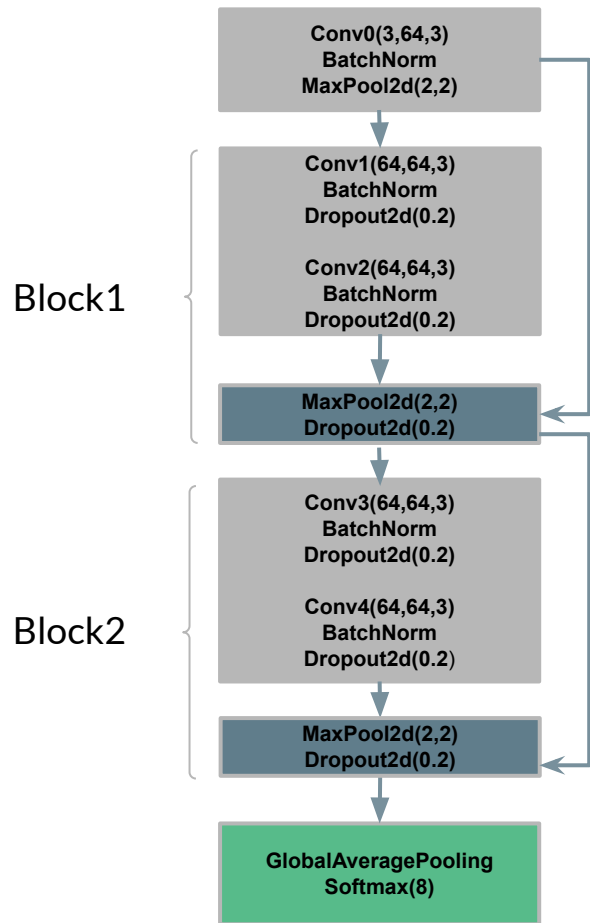
## PyTorch



- **Loss** is much **more stable** on **PyTorch** than on Keras
- Model **performance is preserved**, with very low overfitting in both cases
- For these 100 epochs, **PyTorch had a smaller running time**

# Second Model

# M3 - Model II



## Hyperparameters

Img. Size	64x64
Epochs	500
Batch Size	32
Optimizer	Adam*
Learning Rate	1e-3

### \*Adam configuration


- $lr=0.001$
- $betas=(0.9, 0.999)$
- $eps=1e-07$
- $weight\_decay=0$
- $amsgrad=False$

Activation	ReLu
Weight Init	Glorot uniform
Dropout rate	0.2


# M3 - Model II - Results graphs

## Keras

epoch\_accuracy




epoch\_loss




Training ■  
Validation ■

## PyTorch

Accuracy



Loss



- Validation **accuracy** is a bit **higher in keras** (85%) than in pytorch (80%).
- **Overfitting** is also slightly **worse in pytorch**
- **Testing accuracy is 88%** in both cases, so we might have a mistake on the validation accuracy calculation
  - As opposed to testing, the validation set is obtained as a random split of the training set. So even specifying `shuffle = False`, data changes between experiments, unabling repeatability.
- **Accuracy** seems a bit **more stable in pytorch** and the **loss** decreases on a **nicer** manner (keras drastically drops at the beginning and then plateaus, while PyTorch follows a more reasonable pattern)




# M3 Model II - Comparison II

	Trainable Parameters	Time	Train accuracy	Val accuracy	Test accuracy
Keras	150,664	1805"	98%	85%	88%
PyTorch	150,664	1701"	95%	80%	88%

- We have exactly the same number of trainable parameters in both libraries, which is a good sign
- PyTorch is around 100" faster on these 500 epochs, which is a negligible improvement. Nevertheless, this might become more significant when training bigger models for a larger number of epochs.
- Keras training accuracy is a 3% higher than PyTorch accuracy. This difference is even higher in validation (5%). However, test accuracy is preserved.

Additional investigations

# Data Augmentation I



- We additionally investigated on how to perform data augmentation in PyTorch, as it can be a valuable tool in the future


```
transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((224,224)),
    torchvision.transforms.ColorJitter(brightness=
0.5,contrast=0.6,hue=.05, saturation=.05),
    torchvision.transforms.RandomHorizontalFlip(),
    torchvision.transforms.RandomVerticalFlip(),
    torchvision.transforms.RandomCrop(224),
    torchvision.transforms.RandomRotation(0,
resample=PIL.Image.BILINEAR)
])
```

```
dataset = torchvision.datasets.ImageFolder(images_folder,
transform=transforms)
```

# Data Augmentation II


Probe architecture for data augmentation:

```
class Net(nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.conv1 = nn.Conv2d(3, 6, 5)  
        self.pool = nn.MaxPool2d(2, 2)  
        self.conv2 = nn.Conv2d(6, 16, 5)  
        self.fc1 = nn.Linear(16 * 5 * 5, 120)  
        self.fc2 = nn.Linear(120, 84)  
        self.fc3 = nn.Linear(84, 10)  
  
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = x.view(-1, 16 * 5 * 5)  
        #x = x.view(x.size(0), 16*224*224)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return x
```



Accuracy of the network on the 1881 test images: 71 %  
epoch = 100

# Early stopping




- Early stopping is a form of regularization used to avoid overfitting on the training dataset. It keeps track of the validation loss, if the loss stops decreasing for several epochs in a row the training stops.
- While not implemented by default on PyTorch, there are many community implementations such as [this one](#).
- We set the patience argument in the EarlyStopping class to how many epochs we want to wait after the last time the validation loss improved before breaking the training loop

# Docker image

- Though not directly related to the project, we have taken this opportunity to investigate on how to use Docker in order to have a more versatile application
- This was more relevant when using Keras, as we had some trouble in the past with finding the best configuration of Cuda + Cudnn + Tensorflow + Keras
- Torch seems to work on GPU more easily, but we still found learning about Docker useful

# [Summary]: Keras vs Tensorflow vs Pytorch

- Our M3 model was translated to PyTorch quite straightforwardly
- PyTorch is slightly less user-friendly than Keras, but seems much more versatile and customizable
- Debugging is also probably easier on PyTorch
- GPU implementation was **significantly** easier with PyTorch than with Keras
- Pytorch's documentation is as good or better than Keras, and it seems to have a very active community.



Keras is most suitable for:

- ✓ Rapid Prototyping
- ✓ Small Dataset
  - Multiple back-end support



TensorFlow is most suitable for:

- Large Dataset
- High Performance
- Functionality
- [Object Detection](#)



PyTorch is most suitable for:

- ✓ Flexibility
- ✓ Short Training Duration
- Debugging capabilities

\*<https://www.edureka.co/blog/keras-vs-tensorflow-vs-pytorch/>



Thank You

